




Bayesian Statistics and Hierarchical Bayesian Modeling for Psychological Science

Lecture 09

Lei Zhang

Social, Cognitive and Affective Neuroscience Unit (SCAN-Unit)
Department of Basic Psychological Research and Research Methods

https://github.com/lei-zhang/BayesCog_Wien

lei.zhang@univie.ac.at
lei-zhang.net
 @lei_zhang_lz



universität
wien
Fakultät für Psychologie

Bayesian warm-up?

STAN PROGRAMMING LANGUAGE II



Why Use Stan?

cognitive model

statistics

computing

vs. BUGS and JAGS

- Time to converge and per effective sample size:
0.5 - ∞ times faster
- Memory usage: 1 - 10%
- Language features
 - variable overwrite: $a = 4$, then $a = 5$
 - formal control flow
 - full support of vectorizing



Krzysztof Sakrejda
@sakrejda

I keep getting asked why people should use [@mcmc_stan](#) so I wrote an answer:



"Selling" Stan
discourse.mc-stan.org

27.03.18, 16:01


NUTS vs. Gibbs and Metropolis

cognitive model

statistics

computing

Hamilton MC (HMC) implements No-U-Turn Sampler (NUTS)



- Two dimensions of highly correlated 250-dim normal
- 1,000,000 draws from Metropolis and Gibbs (thin to 1000)
- **1,000** draws from NUTS; 1000 independent draws

General Properties of Stan Language

cognitive model

statistics

computing

- Whitespace does not matter
- Comments
 - `//`
 - `/* ... */`
- Must use semicolon (;)
- Variables are typed and scoped



Variable's Scope

cognitive model

statistics

computing

	data	transformed data	parameters	transformed parameters	model	generated quantities
Variable Declarations	Yes	Yes	Yes	Yes	Yes	Yes
Variable Scope	Global	Global	Global	Global	Local	Local
Variables Saved?	No	No	Yes	Yes	No	Yes
Modify Posterior?	No	No	No	No	Yes	No
Random Variables	No	No	No	No	No	Yes

Variable Declaration

cognitive model

statistics

computing

- Each variable has a type (static type; scalar, vector, matrix etc.)
- Only values of that type can be assigned to the variable
 - e.g. cannot assign `[1 2 3]` to `a` (declared as a scalar)
- Declaration of variables happen at the top of a block (including local blocks)



Scalar Variables

real

- scalar
- continuous

```
data {  
  real y;  
}
```

int

- scalar
- integer
- can't be used in parameters or transformed parameters blocks

```
data {  
  int n;  
}
```

Constraining Scalar Variables

```
data {  
  int<lower=1> m;  
  int<lower=0,upper=1> n;  
  real<lower=0> x;  
  real<upper=0> y;  
  real<lower=-1,upper=1> rho;  
}
```

Vector & Matrix

```
vector<double> a;  
// column vector
```

```
row_vector<double> b;  
// row vector
```

```
matrix<double> A;  
// A is a 3x4 matrix  
// A[1] returns a 4-element row vector
```

```
vector<double> rhos;  
row_vector<double> sigmas;  
matrix<double> Sigma;
```

Control Flow

- if-else

```
if (cond) {  
  ..statement..  
}
```

```
if (cond) {  
  ..statement..  
} else {  
  ..statement..  
}
```

```
if (cond) {  
  ..statement..  
} else if (cond) {  
  ..statement..  
} else {  
  ..statement..  
}
```

- for-loop

```
for ( j in 1:J ) {  
  ..statement..  
}
```

```
for ( j in 1:J ) {  
  for ( k in 1:K ) {  
    ..statement..  
  }  
}
```

same as the R syntax, but
terminate each line with ;

BERNOULLI MODEL



Bernoulli Model

cognitive model

statistics

computing

- You are interested in if a coin is biased.
- You will flip the coin.
- You will record whether it comes up a head (h) or a tail (t).
- You might observe 15 heads out of 20 flips.
- What is your degree of belief about the biased parameter ϑ ?



Bernoulli Model

cognitive model


statistics

computing

$$p(w | N, p) = \binom{N}{w} p^w (1 - p)^{N-w}$$

$N = 1$

$$p(h | \theta) = \theta^h (1 - \theta)^{1-h}$$



$\theta \sim \text{Uniform}(0, 1)$

$h \sim \text{Bernoulli}(\theta)$

Exercise VIII

cognitive model


statistics

computing

```
.../BayesCog/03.bernoulli_coin/_scripts/bernoulli_coin_main.R
```

TASK: fit the Bernoulli model

```
> dataList  
$`flip`  
[1] 1 1 1 0 1 1 1 1 1 0 0 1 1 0 1 1 1 1 0 1  
  
$N  
[1] 20
```



Possible Optimization?

cognitive model

statistics

computing

```
model {  
  for (n in 1:N) {  
    flip[n] ~ bernoulli(theta);  
  }  
}
```

61.59 secs*

```
model {  
  flip ~ bernoulli(theta);  
}
```

53.25 secs*

Thinking before looping!

Recap

What we've learned...

- R Basics
- probability distributions
- Bayes' theorem, $p(\theta|D)$
- Binomial model
- MCMC and Stan

LINEAR REGRESSION


$y(x_0)$

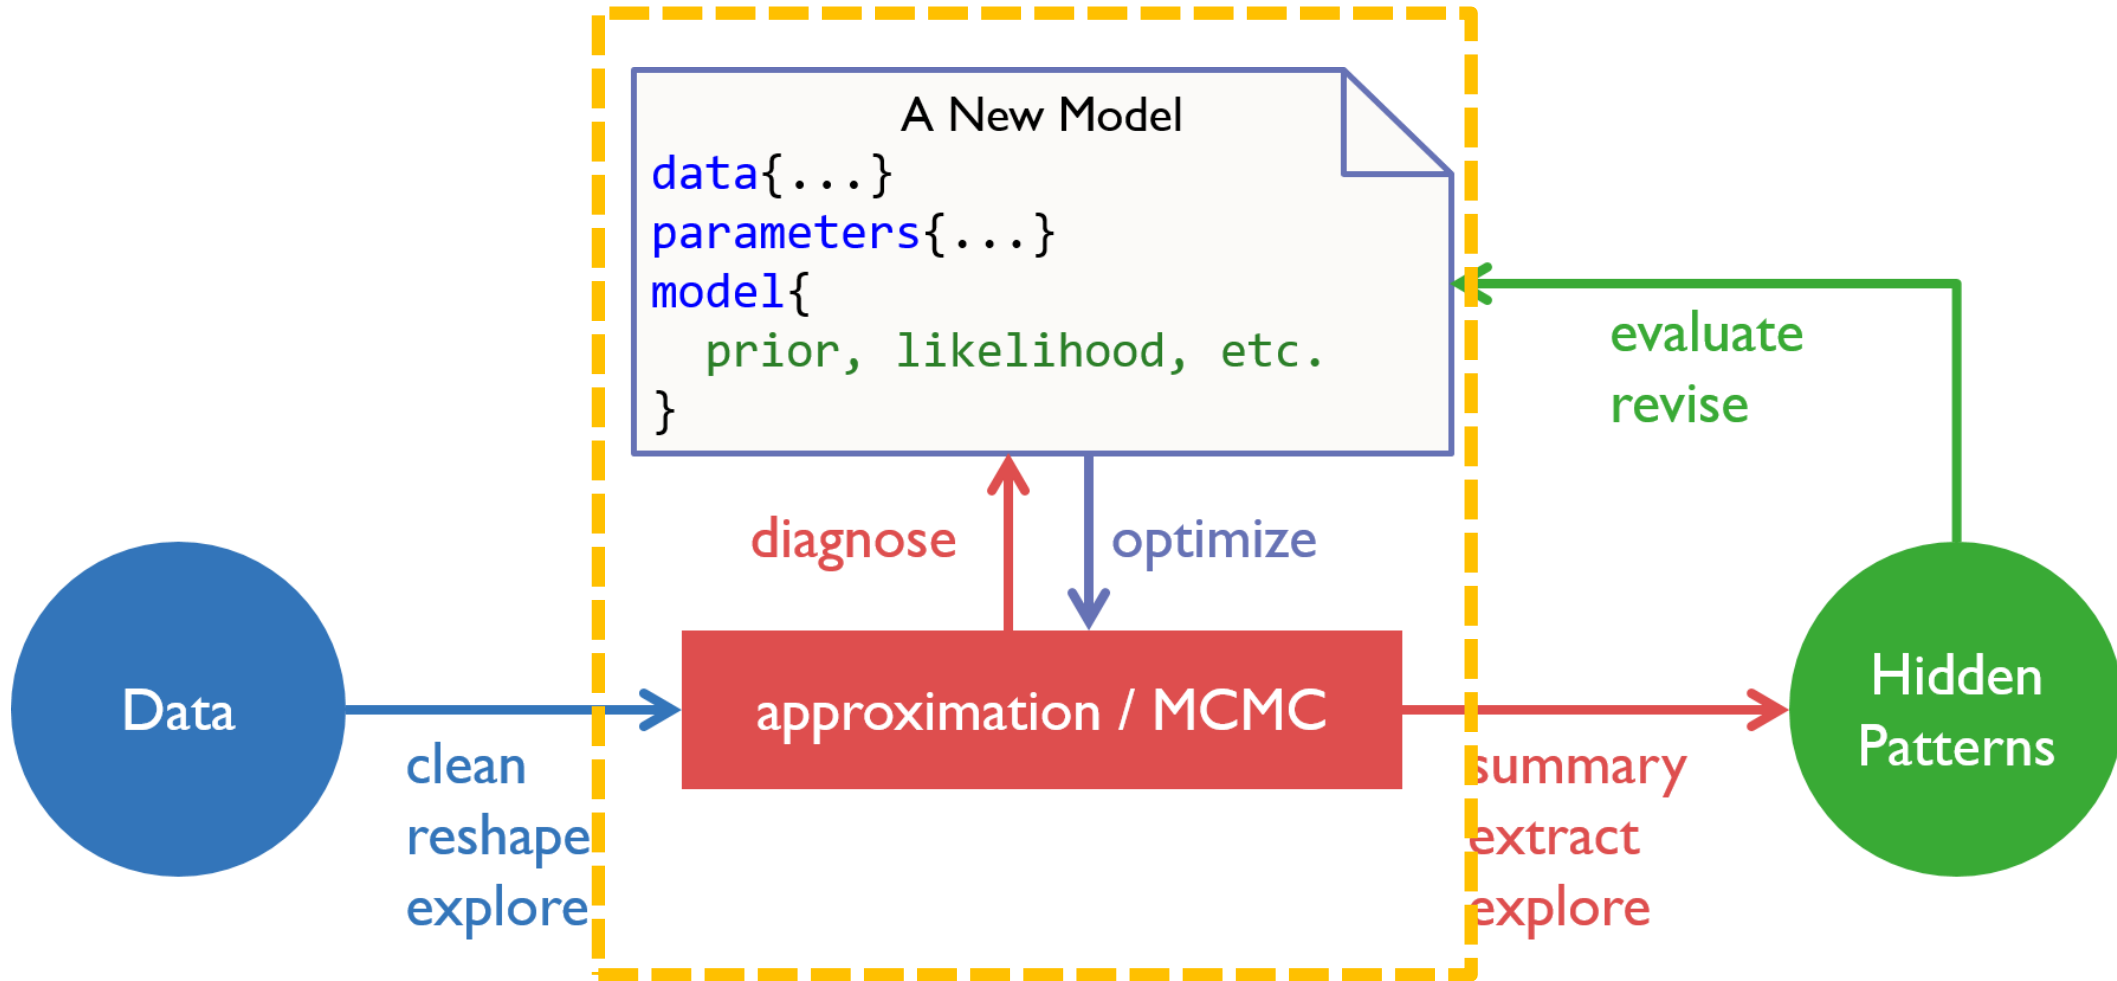
$p(t|x_0)$

$y(x)$

x_0

x





Linear Regression: height ~ weight

cognitive model


statistics

computing

```
.../04.regression_height/_scripts/regression_height_main.R
```

make scatter plot and fit the model with `lm()`

```
>load('_data/height.RData')
>d <- Howell1
>d <- d[ d$age >= 18 , ]
>head(d)
height  weight age male
1 151.765 47.82561 63    1
2 139.700 36.48581 63    0
3 136.525 31.86484 65    0
4 156.845 53.04191 41    1
5 145.415 41.27687 51    0
6 163.830 62.99259 35    1
```



Results with lm()

cognitive model

statistics

computing

```
> L <- lm( height ~ weight, d) # estimate model by minimizing least squares errors
> summary(L)
```

Call:

```
lm(formula = height ~ weight, data = d)
```

Residuals:

Min	1Q	Median	3Q	Max
-19.7464	-2.8835	0.0222	3.1424	14.7744

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	113.87939	1.91107	59.59	<2e-16	***
weight	0.90503	0.04205	21.52	<2e-16	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 5.086 on 350 degrees of freedom

Multiple R-squared: 0.5696, Adjusted R-squared: 0.5684


F-statistic: 463.3 on 1 and 350 DF, p-value: < 2.2e-16

height ~ weight

cognitive model

statistics

computing



Rethinking Regression Model

cognitive model

statistics

computing

$$\mu_i = \alpha + \beta x_i$$

~~$$y_i = \mu_i + \varepsilon$$~~

~~$$\varepsilon \sim \text{Normal}(0, \sigma)$$~~

$$y_i \sim \text{Normal}(\mu_i, \sigma)$$


Rethinking Regression Model

cognitive model

statistics

computing

$$\mu_i = \alpha + \beta x_i$$
$$y_i \sim \text{Normal}(\mu_i, \sigma)$$




Rethinking Regression Model

cognitive model

statistics

computing

$$\mu_i = \alpha + \beta x_i$$
$$y_i \sim \text{Normal}(\mu_i, \sigma)$$



```
model {  
  vector[N] mu;  
  for (i in 1:N) {  
    mu[i] = alpha + beta * weight[i];  
    height[i] ~ normal(mu[i], sigma);  
  }  
}
```

```
model {  
  vector[N] mu;  
  mu = alpha + beta * weight;  
  height ~ normal(mu, sigma);  
}
```

```
model {  
  height ~ normal(alpha + beta * weight, sigma);  
}
```

Thinking about Priors?

cognitive model

statistics


computing

$\alpha \sim \text{Normal}(170, 100)$ $\beta \sim \text{Normal}(0, 20)$

$\overline{\text{height}} = \alpha + \beta * \text{weight}$

$\sigma \sim \text{halfCauchy}(0, 20)$

$\text{height} \sim \text{Normal}(\overline{\text{height}}, \sigma)$



Exercise VIII

cognitive model

statistics

computing

```
.../04.regression_height/_scripts/regression_height_main.R
```

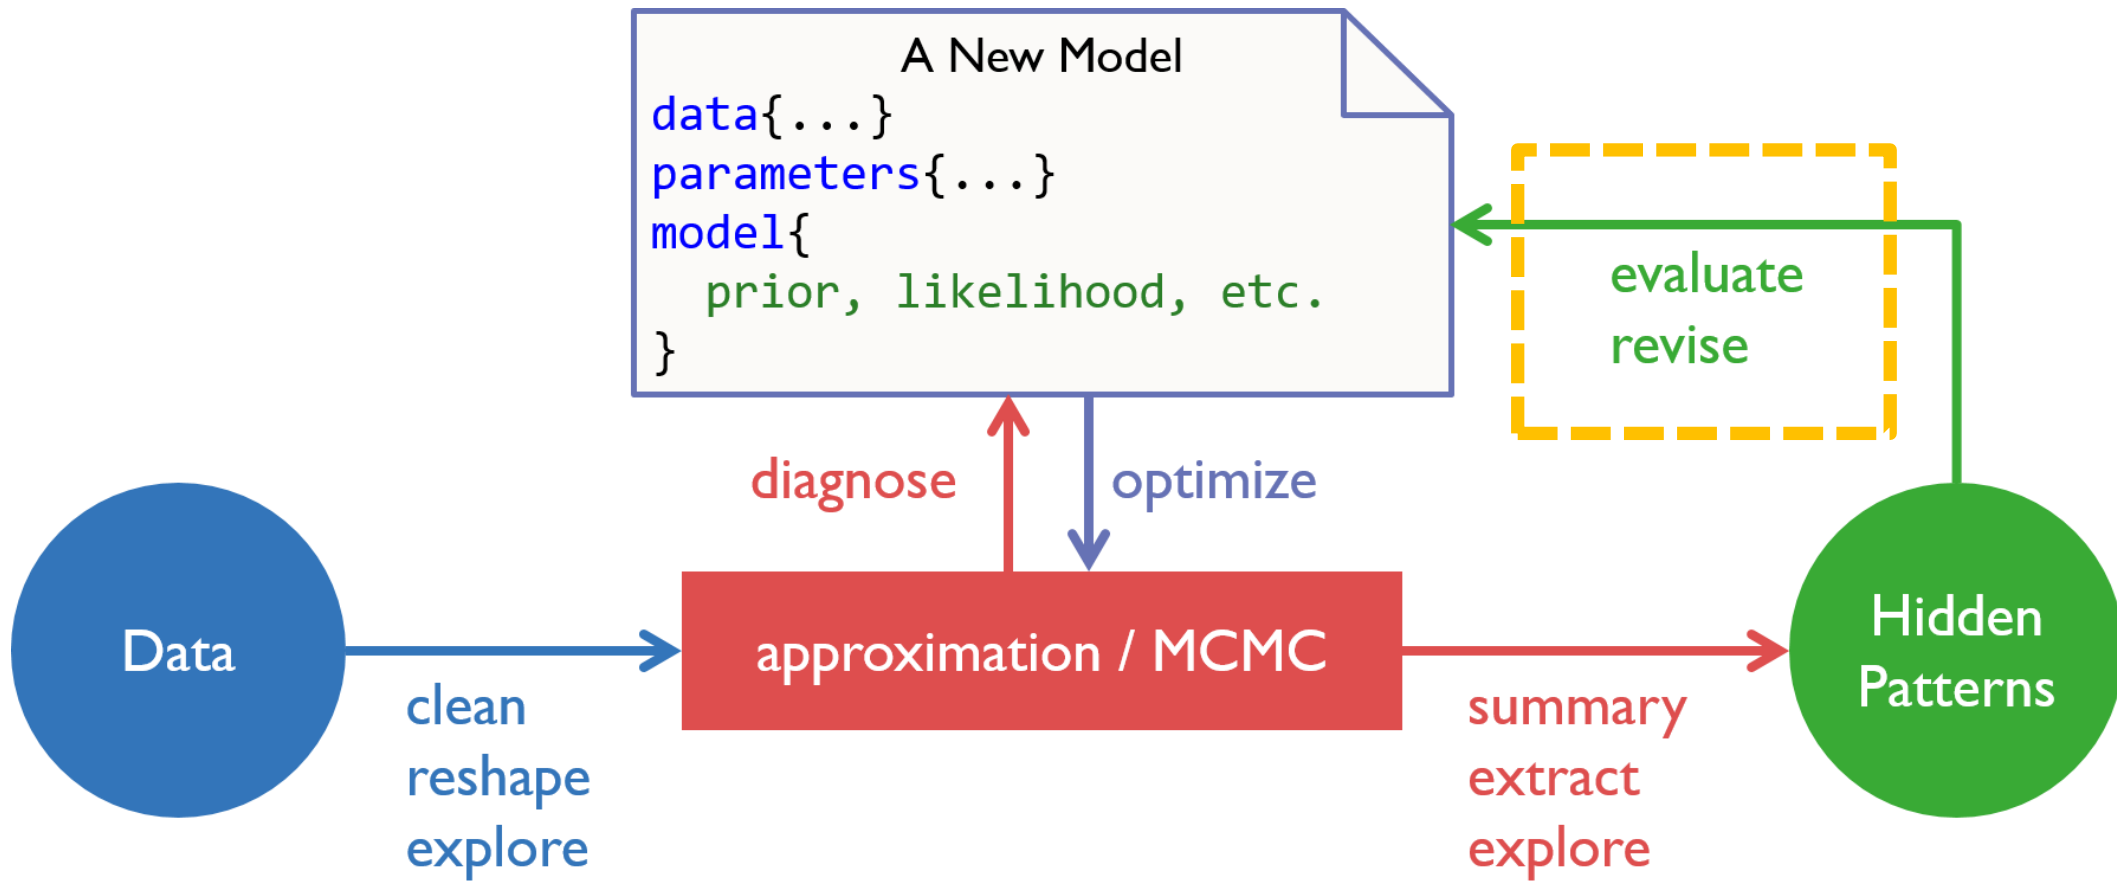
TASK: estimate the model and produce the results

Inference for Stan model: regression_height_model.

4 chains, each with iter=2000; warmup=1000; thin=1;

post-warmup draws per chain=1000, total post-warmup draws=4000.

	mean	se_mean	sd	2.5%	25%	50%	75%	97.5%	n_eff	Rhat
alpha	113.97	0.06	1.86	110.27	112.76	113.93	115.20	117.66	934	1
beta	0.90	0.00	0.04	0.82	0.88	0.90	0.93	0.99	922	1
sigma	5.11	0.01	0.19	4.74	4.97	5.10	5.24	5.50	1437	1
lp__	-747.61	0.04	1.23	-750.80	-748.15	-747.28	-746.72	-746.24	993	1




What does the Model Predict?

cognitive model

statistics

computing



$$p(y_{rep} | y) = \int p(y_{rep} | \theta) p(\theta | y) d(\theta)$$

Posterior Predictive Check (PPC)


cognitive model

statistics

computing

```
generated quantities {  
  vector[N] height_bar;  
  for (n in 1:N) {  
    height_bar[n] = normal_rng(alpha + beta * weight[n], sigma);  
  }  
}
```

the generated quantities block runs only AFTER the sampling, and the time it costs can be essentially ignored!



Posterior Predictive Check (PPC)

cognitive model



statistics

computing

Highest density interval
(HDI)

`dens(height_bar | x=47.8)`

possible HDI values: 80%, 89%, 90%, 95%




```
height_bar <- extract(fit_reg_ppc, pars = 'height_bar',  
                      permuted = FALSE)$height_bar  
height_HDI <- apply(height_bar, 2, HDIoFMC)
```


Posterior Predictive Check (PPC)

cognitive model

statistics

computing



Exercise IX


cognitive model

statistics

computing

```
.../05.regression_height_poly/_scripts  
/regression_height_poly_main.R
```

- TASK: (1) Complete “regression_height_poly2_model.stan”
(2) produce PPC plot for both 1st order and 2nd order polynomial fit



Exercise IX – Tips

cognitive model

statistics

computing

```
> source('_scripts/regression_height_poly_main.R')
```

```
> out1 <- reg_poly(poly_order = 1)
```

$$\overline{\text{height}} = \alpha + \beta_1 * \text{weight} + \beta_2 * \text{weight}^2$$

$$\text{height} \sim \text{Normal}(\overline{\text{height}}, \sigma)$$

```
data {  
  int<lower=0> N;  
  vector<lower=0>[N] height;  
  vector<lower=0>[N] weight;  
  vector<lower=0>[N] weight_sq;  
}
```

```
height ~ normal(alpha + beta1 * weight + beta2 * weight_sq, sigma);
```

Exercise IX – output2

cognitive model

statistics

computing

