




Bayesian Statistics and Hierarchical Bayesian Modeling for Psychological Science

Lecture 12

Lei Zhang

Social, Cognitive and Affective Neuroscience Unit (SCAN-Unit)
Department of Basic Psychological Research and Research Methods

https://github.com/lei-zhang/BayesCog_Wien

lei.zhang@univie.ac.at
lei-zhang.net
 @lei_zhang_lz



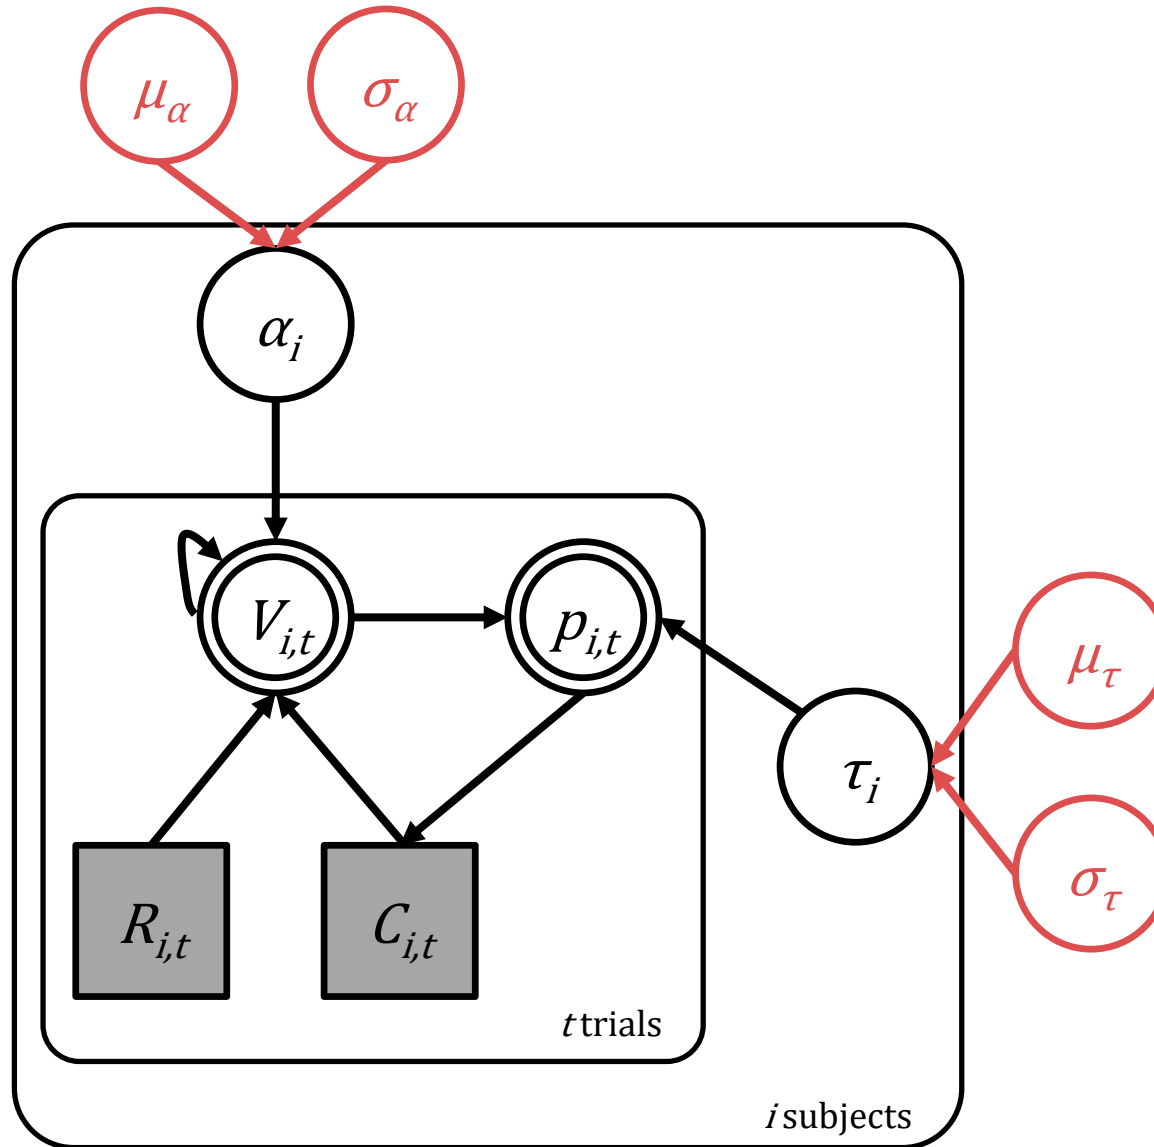
universität
wien
Fakultät für Psychologie

Hierarchical RL Model

cognitive model

statistics

computing

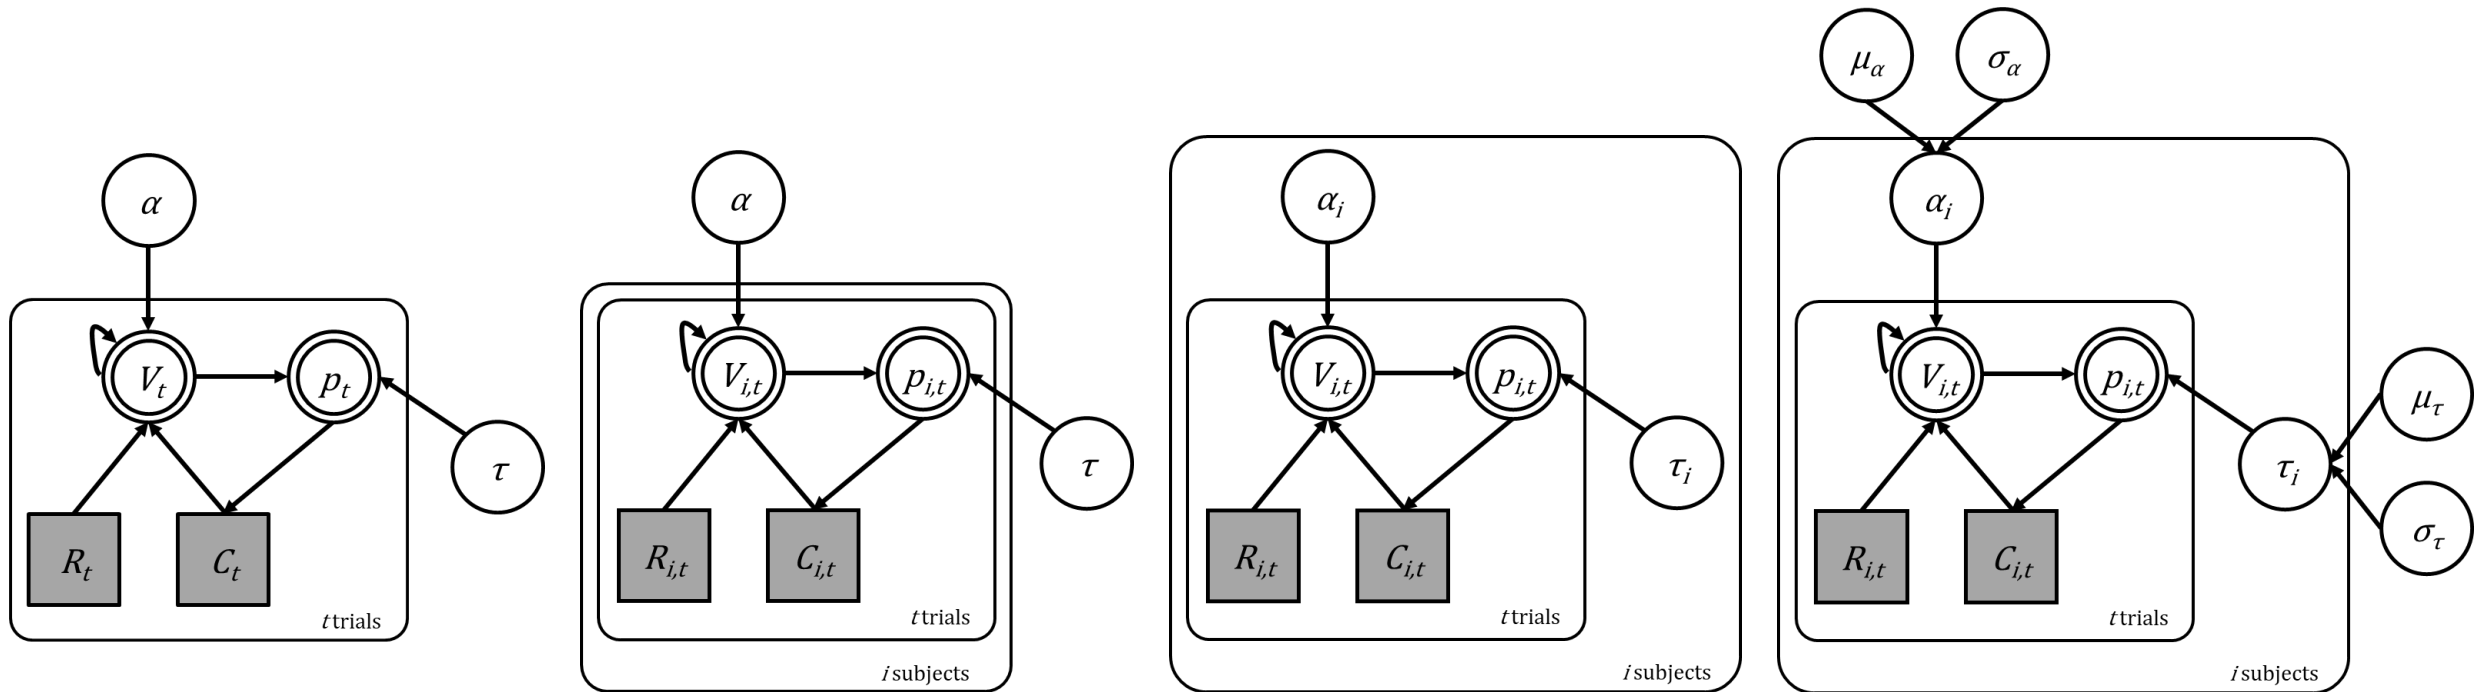


HOW DID WE GET HERE?

cognitive model

statistics

computing



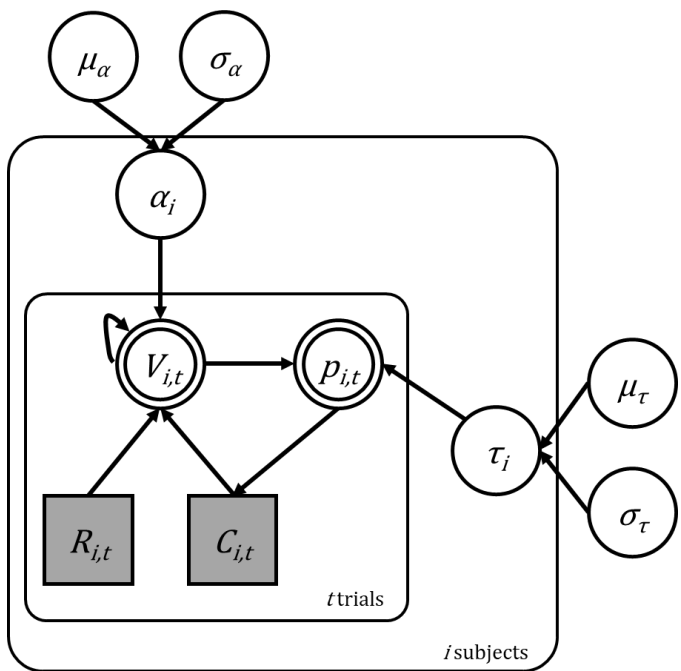
The cognitive model *per se* is the same!

Implementing Hierarchical RL Model

cognitive model

statistics

computing



$$\mu_{\alpha} \sim \text{Uniform}(0, 1)$$

$$\sigma_{\alpha} \sim \text{halfCauchy}(0, 1)$$

$$\mu_{\tau} \sim \text{Uniform}(0, 3)$$

$$\sigma_{\tau} \sim \text{halfCauchy}(0, 3)$$

$$\alpha_i \sim \text{Normal}(\mu_{\alpha}, \sigma_{\alpha}) \tau(0, 1)$$

$$\tau_i \sim \text{Normal}(\mu_{\tau}, \sigma_{\tau}) \tau(0, 3)$$

$$p_{i,t}(C = A) = \frac{1}{1 + e^{\tau_i(V_{i,t}(B) - V_{i,t}(A))}}$$

$$V_{i,t+1}^C = V_{i,t}^C + \alpha_i(R_{i,t} - V_{i,t}^C)$$

```
parameters {
  real<lower=0,upper=1> lr_mu;
  real<lower=0,upper=3> tau_mu;

  real<lower=0> lr_sd;
  real<lower=0> tau_sd;

  real<lower=0,upper=1> lr[nSubjects];
  real<lower=0,upper=3> tau[nSubjects];
}
```

```
model {
  lr_sd ~ cauchy(0, 1);
  tau_sd ~ cauchy(0, 3);
  lr ~ normal(lr_mu, lr_sd) ;
  tau ~ normal(tau_mu, tau_sd) ;
}
```

```
for (s in 1:nSubjects) {
  vector[2] v;
  real pe;
  v = initV;

  for (t in 1:nTrials) {
    choice[s,t] ~ categorical_logit( tau[s] * v );
    pe = reward[s,t] - v[choice[s,t]];
    v[choice[s,t]] = v[choice[s,t]] + lr[s] * pe;
  }
}
```

Exercise XI

cognitive model

statistics

computing

```
.../06.reinforcement_learning/_scripts/reinforcement_learning_multi_parm_main.R
```

TASK: (1) complete the model (TIP: individual ~ group)
(2) fit the hierarchical RL model

```
> source('_scripts/reinforcement_learning_multi_parm_main.R')  
  
> fit_rl3 <- run_rl_mp( modelType = 'hrch' )
```

In addition: Warning messages:

1: There were 97 divergent transitions after warmup. Increasing adapt_delta above 0.8 may help. See <http://mc-stan.org/misc/warnings.html#divergent-transitions-after-warmup>

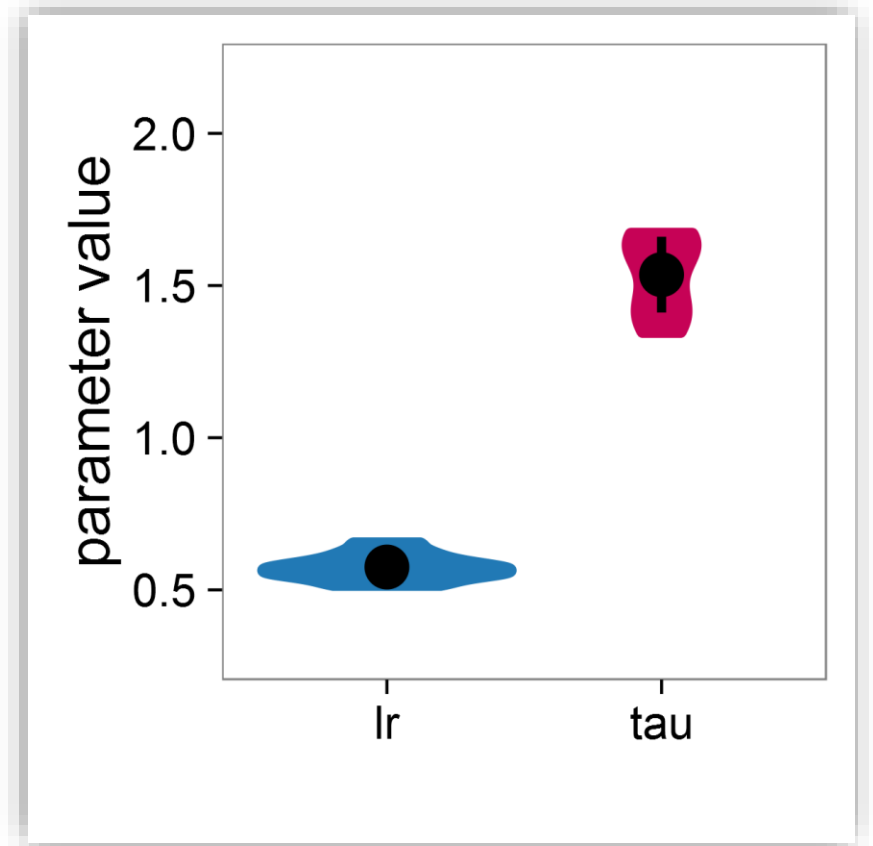
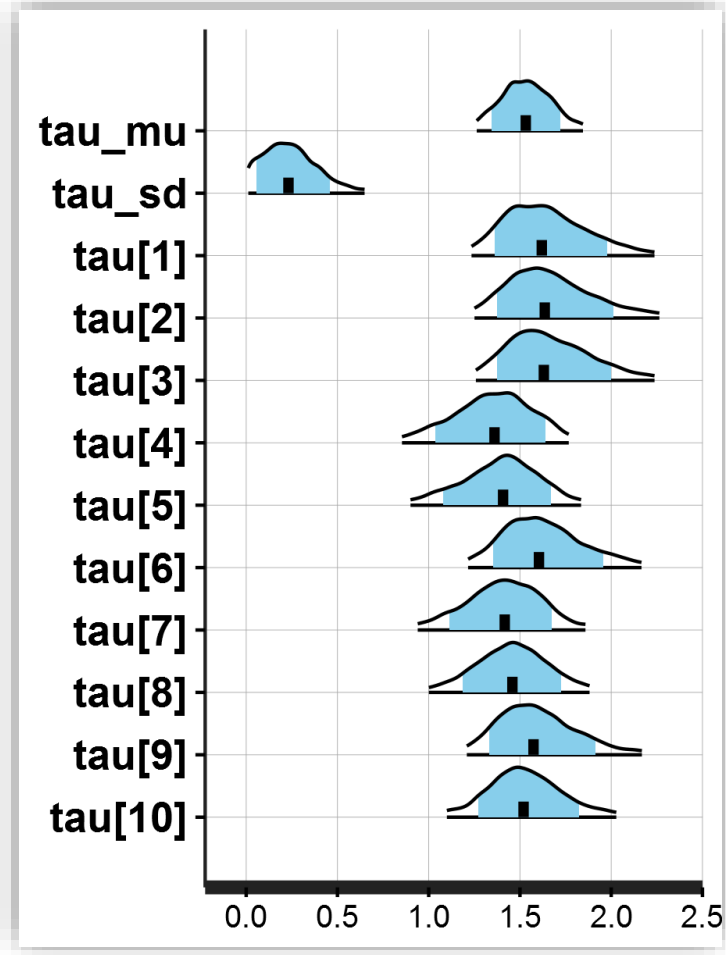
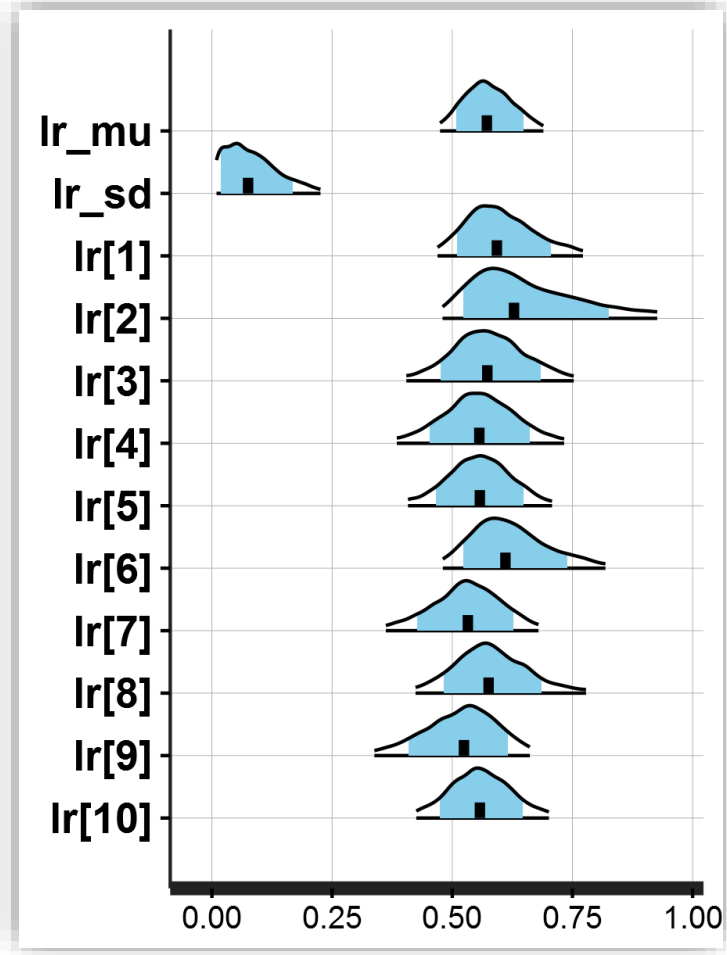
2: Examine the pairs() plot to diagnose sampling problems

Hierarchical Fitting*

cognitive model

statistics

computing



*: adapt_delta=0.999, max_treedepth=100

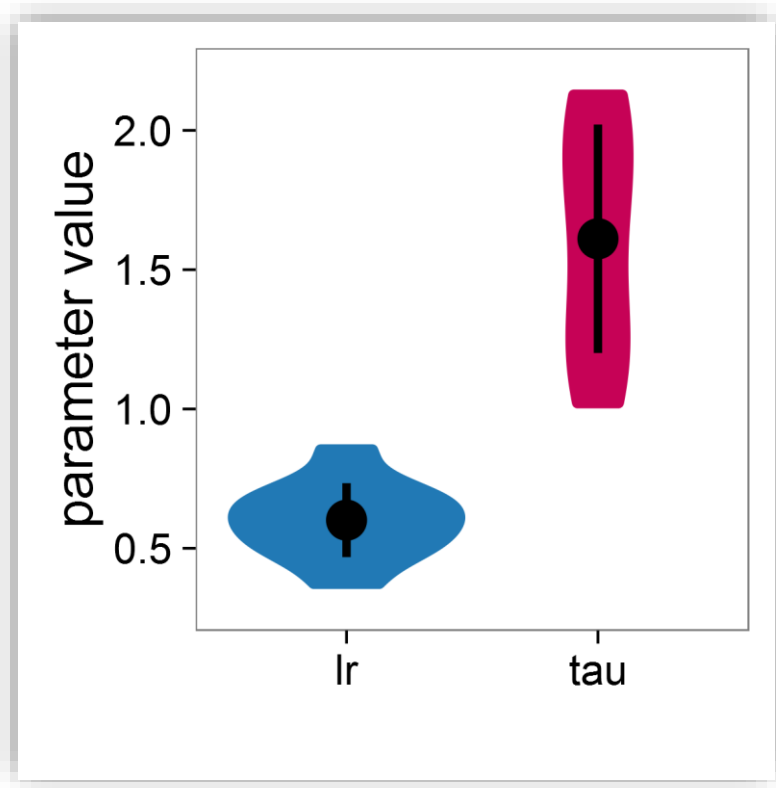
Comparing with True Parameters

cognitive model

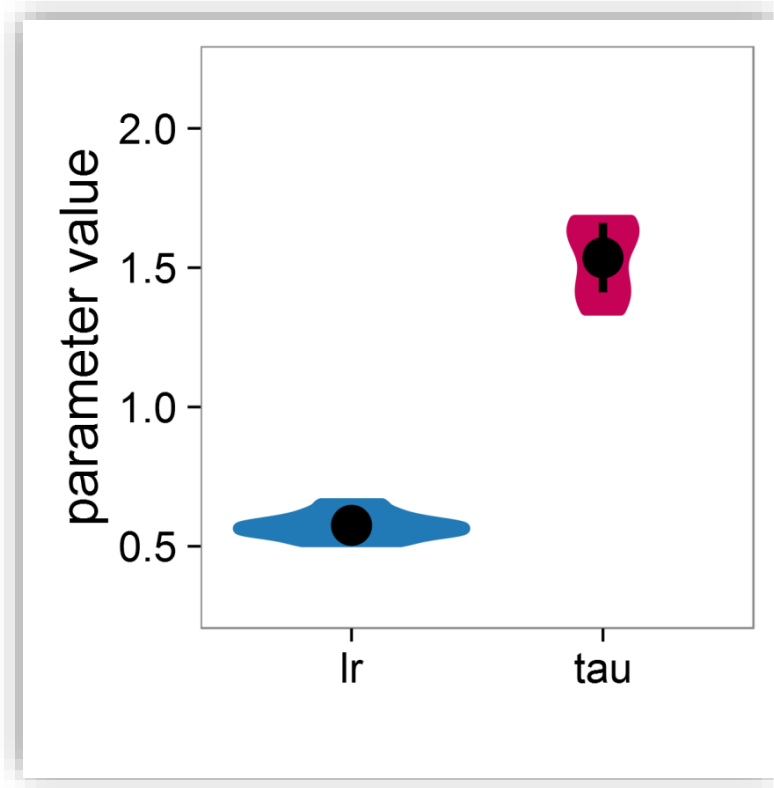
statistics

computing

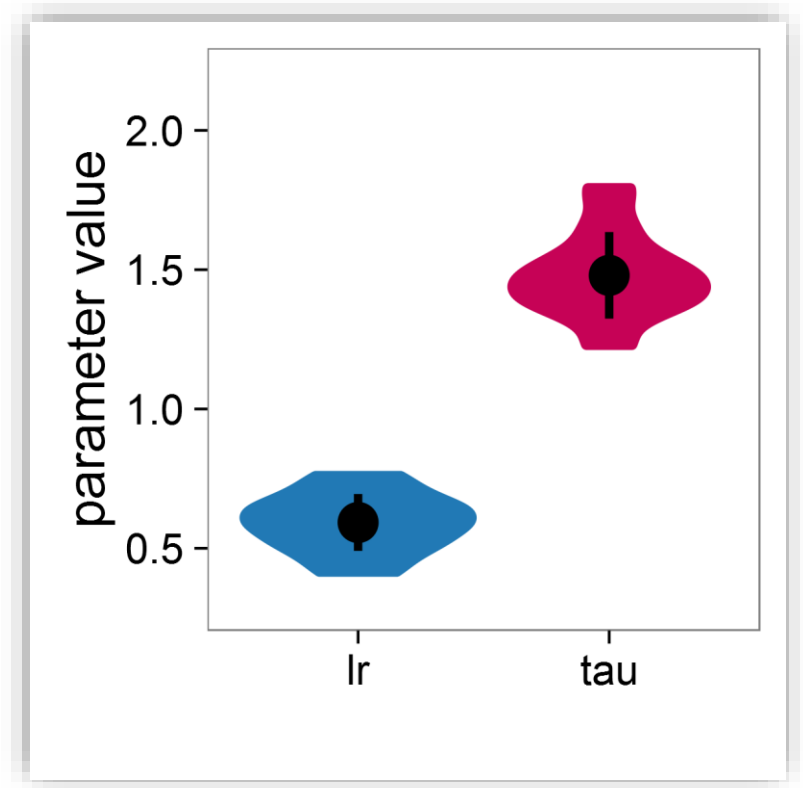
Posterior Means (indv)



Posterior Means (hrch)*



True Parameters



*: adapt_delta=0.999, max_treedepth=100

Group-level Parameters

cognitive model

statistics

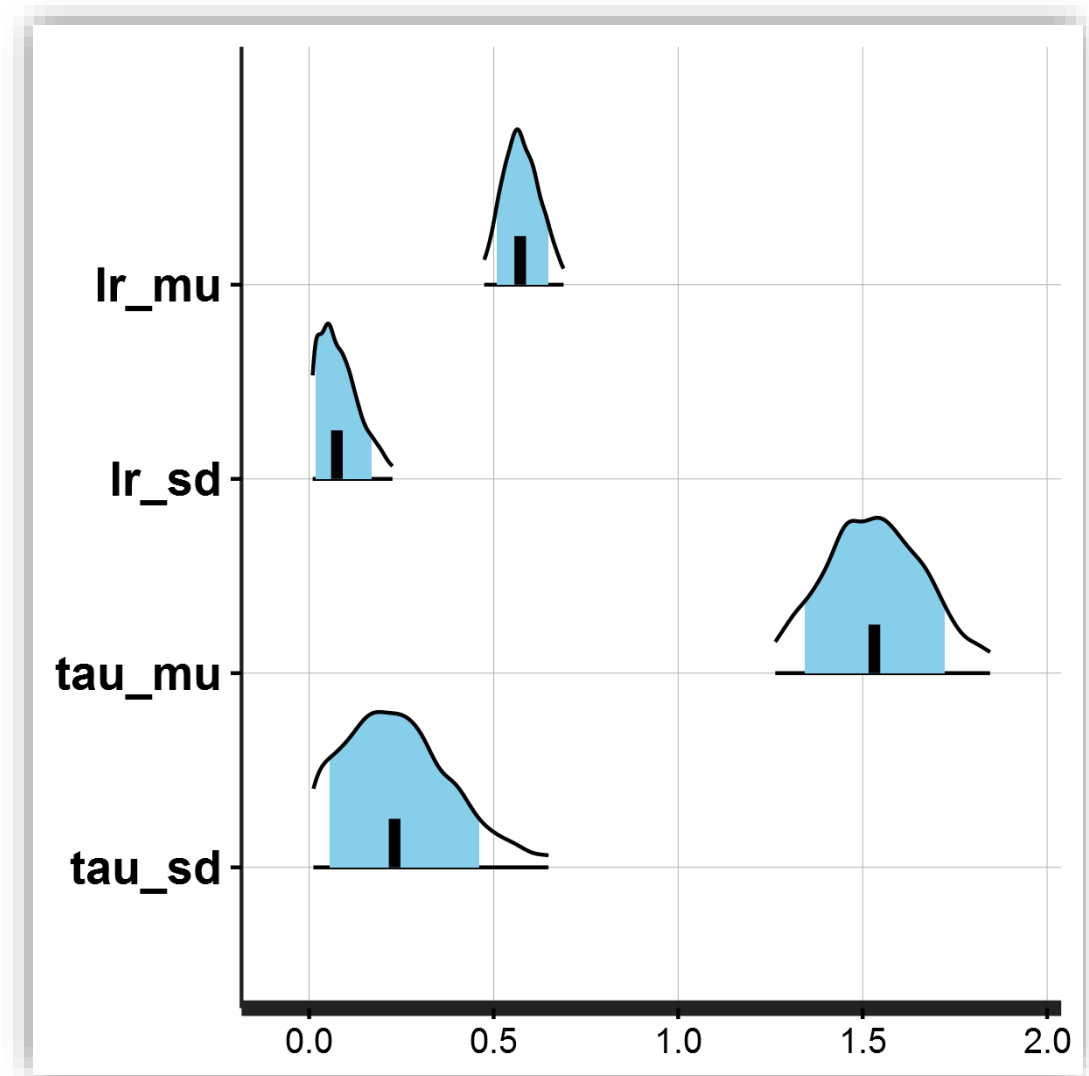
computing

True group parameters

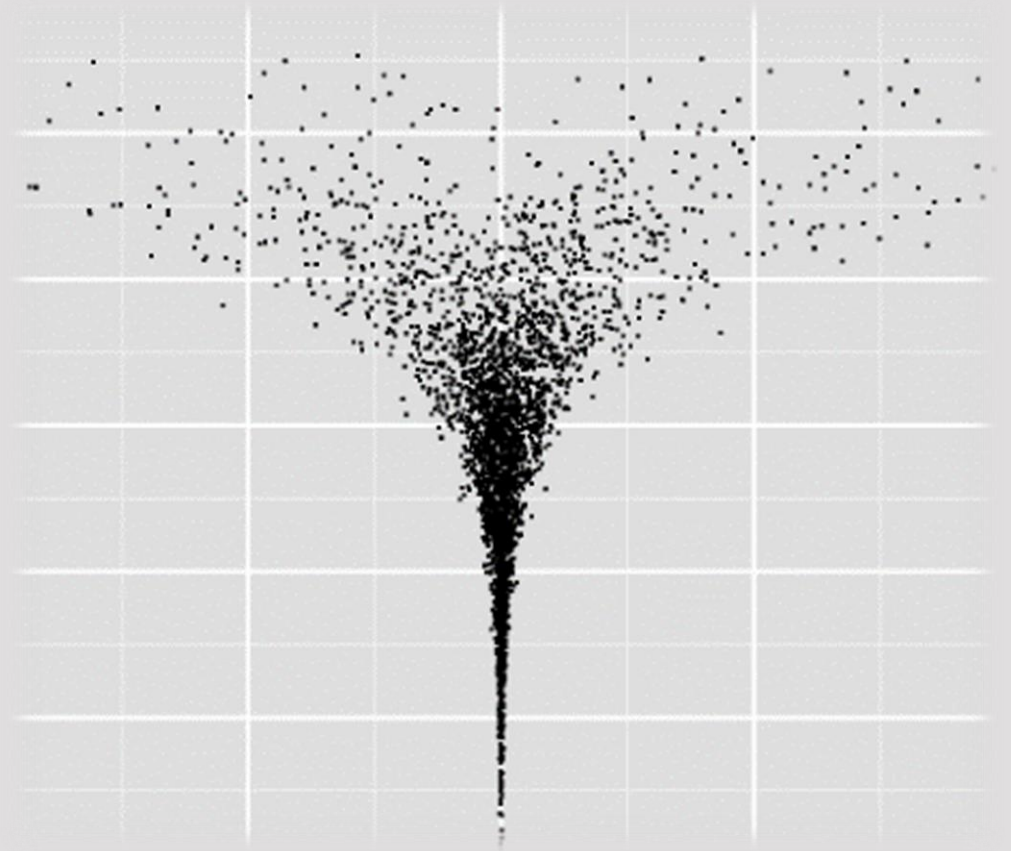
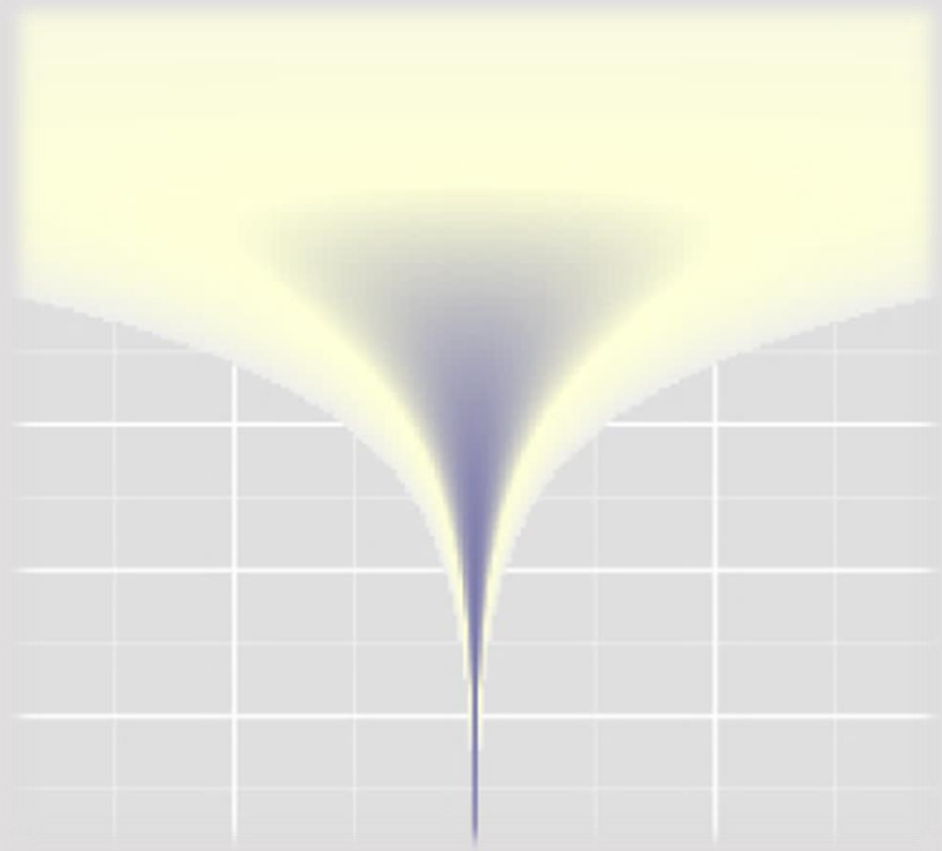
```
lr = rnorm(10, mean=0.6, sd=0.12)
tau = rnorm(10, mean=1.5, sd=0.2)
```

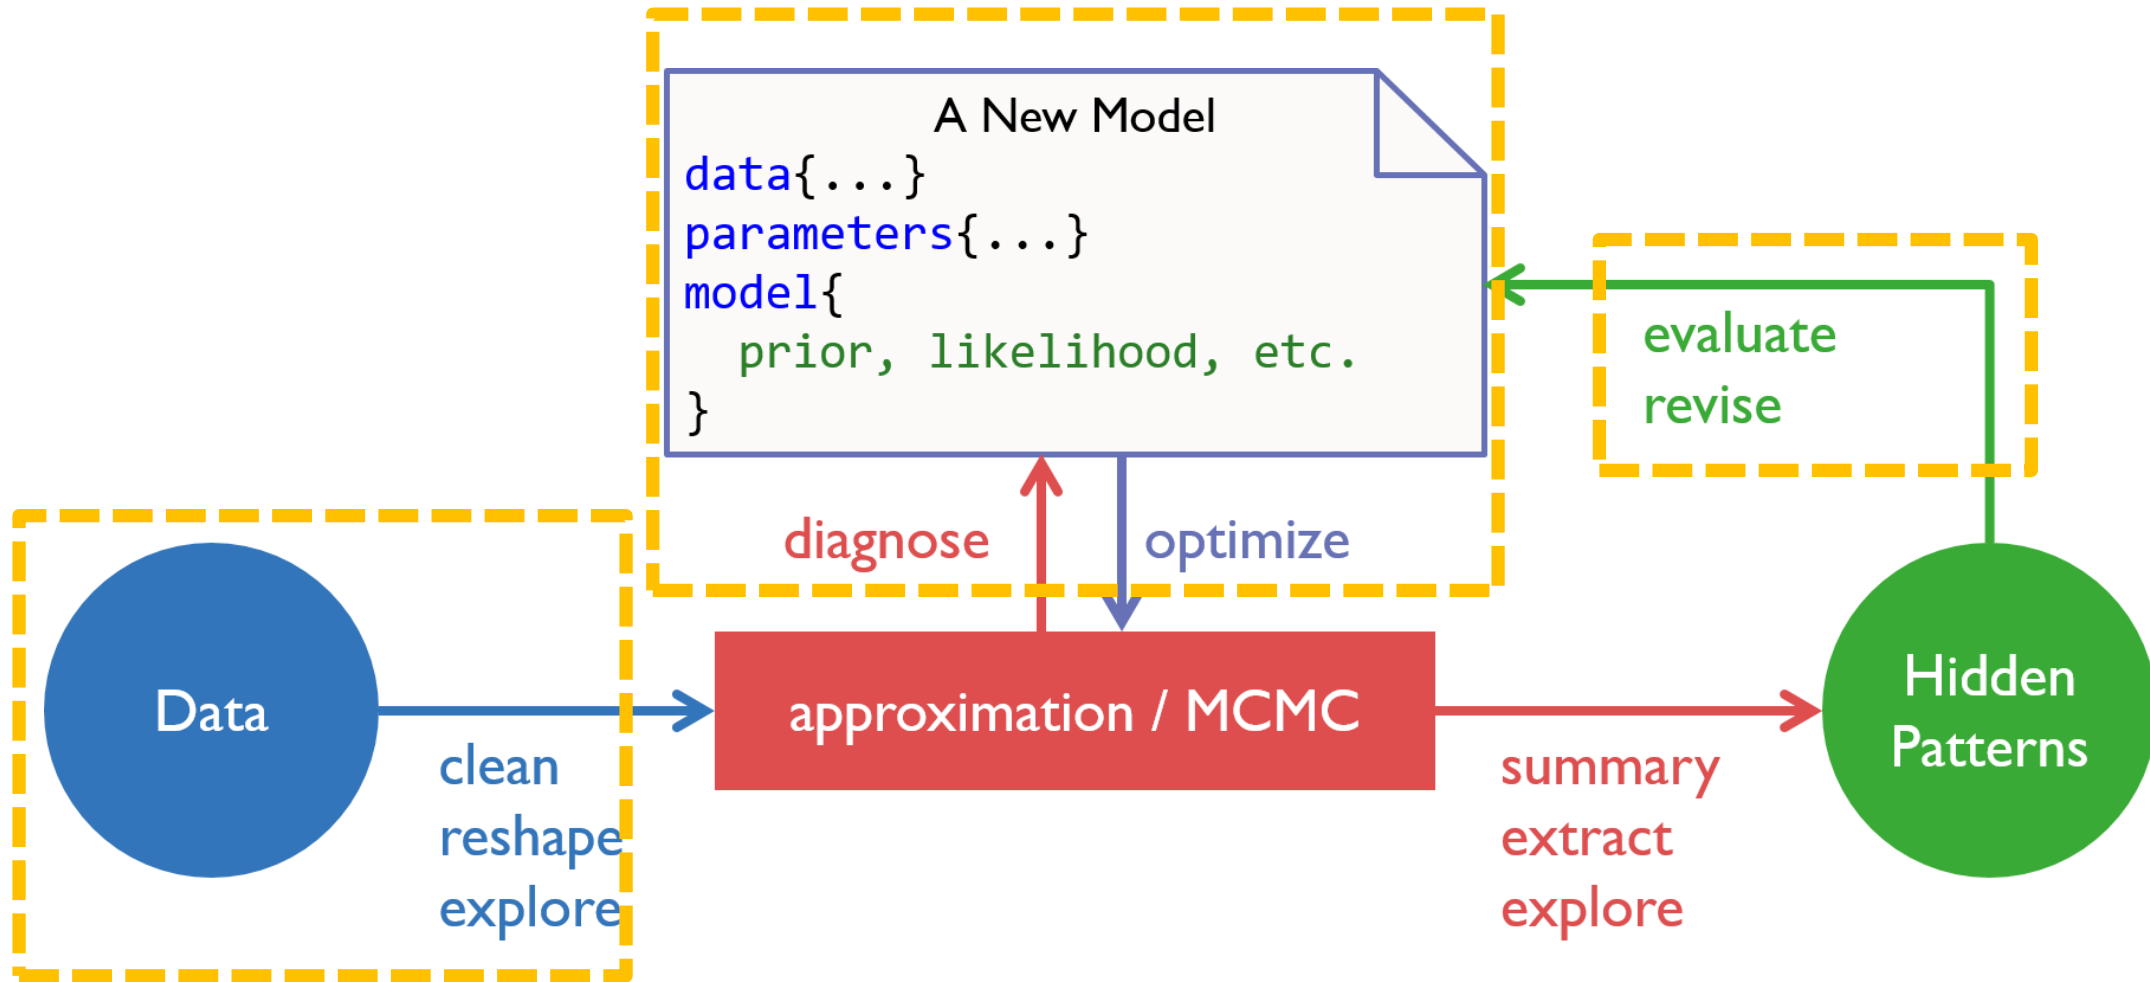
Estimated group parameters

	mean	2.5%	25%	50%	75%	97.5%
lr_mu	0.58	0.47	0.54	0.57	0.61	0.69
lr_sd	0.09	0.01	0.04	0.08	0.12	0.23
tau_mu	1.54	1.26	1.43	1.53	1.63	1.85
tau_sd	0.25	0.01	0.13	0.23	0.34	0.65



OPTIMIZING STAN CODES







Optimizing Stan Code

cognitive model

statistics

computing

Preprocess data

run as many calculations as you can outside Stan

Specify a proper model

follow literature, supervision, experience, etc.

Vectorizing

vectorize Stan code whenever you can

Reparameterizing

reparameterize target parameter to simple distributions

Preprocess Data

cognitive model

statistics

computing

$$\overline{\text{height}} = \alpha + \beta_1 * \text{weight} + \beta_2 * \text{weight}^2$$

```
d$weight_sq <- d$weight^2
```

```
data {  
  int<lower=0> N;  
  vector<lower=0>[N] height;  
  vector<lower=0>[N] weight;  
  vector<lower=0>[N] weight_sq;  
}
```

Specify a Proper Model

cognitive model

statistics

computing

- Visualize your data
- Follow Literatures
- Start from simple and then build complexities
- Simulate data and run model recovery

A New Model

```
data{...}  
parameters{...}  
model{  
  prior, likelihood, etc.  
}
```

Vectorization

cognitive model

statistics

computing

```
model {  
  for (n in 1:N) {  
    flip[n] ~ bernoulli(theta);  
  }  
}
```



```
model {  
  flip ~ bernoulli(theta);  
}
```

```
parameters {  
  ...  
  real<lower=0,upper=1> lr[nSubjects];  
  real<lower=0,upper=3> tau[nSubjects];  
}  
  
model {  
  ...  
  lr ~ normal(lr_mu, lr_sd) ;  
  tau ~ normal(tau_mu, tau_sd) ;  
  ...  
}
```

```
model {  
  vector[N] mu;  
  for (i in 1:N) {  
    mu[i] = alpha + beta * weight[i];  
    height[i] ~ normal(mu[i], sigma)  
  }  
}
```



```
model {  
  vector[N] mu;  
  mu = alpha + beta * weight;  
  height ~ normal(mu, sigma);  
}
```



```
model {  
  height ~ normal(alpha + beta * weight, sigma);  
}
```

Reparameterization

Neal's Funnel

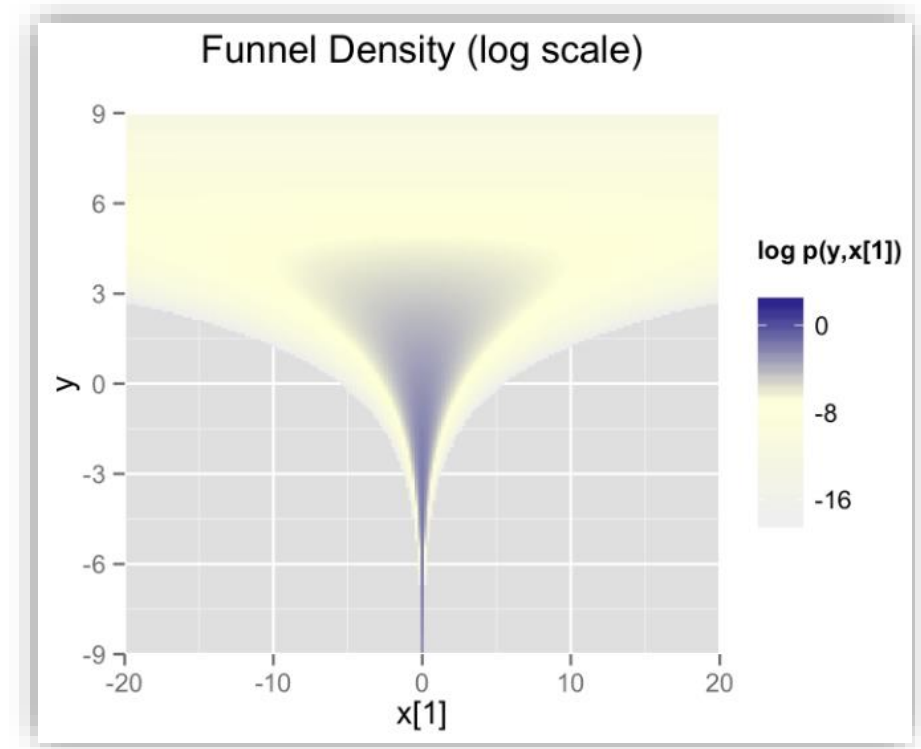
cognitive model

statistics

computing

$$p(y, x) = \text{Normal}(y|0, 3) \times \prod_{n=1}^9 \text{Normal}(x_n|0, \exp(y/2))$$

```
parameters {  
  real y;  
  vector[9] x;  
}  
model {  
  y ~ normal(0, 3);  
  x ~ normal(0, exp(y/2));  
}
```



Non-centered Reparameterization*

cognitive model

statistics

computing

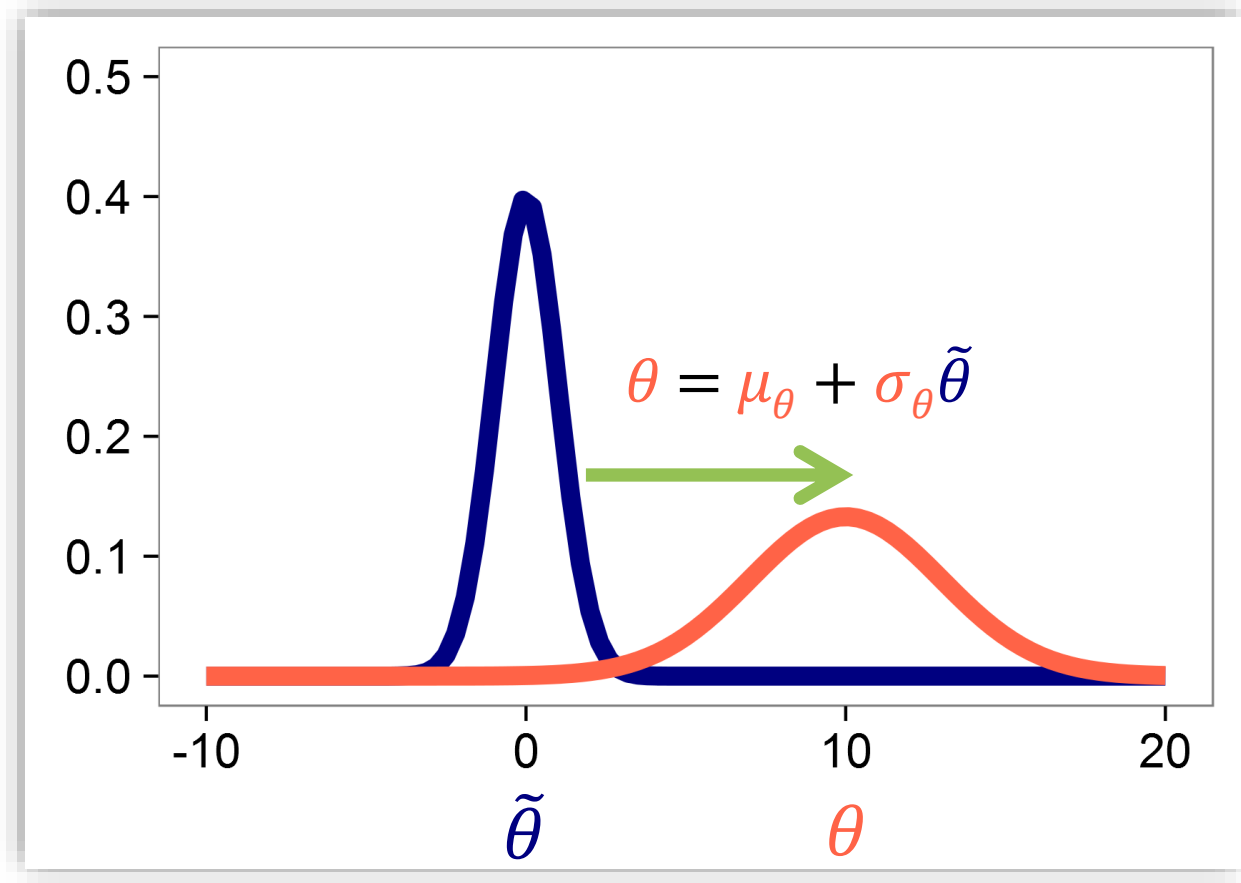
$$\theta \sim \text{Normal}(\mu_\theta, \sigma_\theta)$$



$$\tilde{\theta} \sim \text{Normal}(0, 1)$$

$$\theta = \mu_\theta + \sigma_\theta \tilde{\theta}$$

Stan likes **simple**
distributions!



Reparameterization

Neal's Funnel

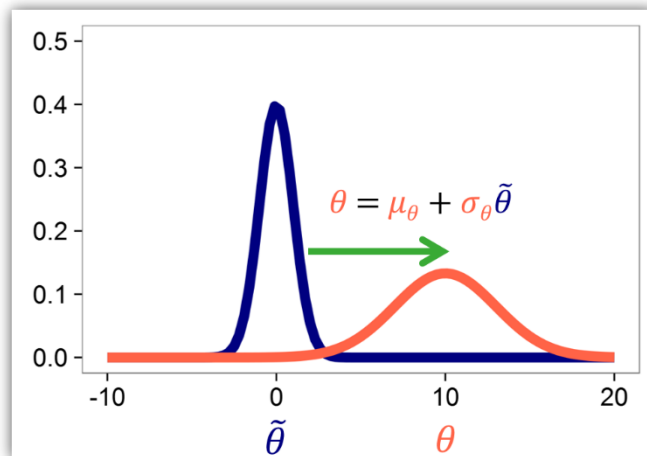
$$p(y, x) = \text{Normal}(y|0, 3) \times \prod_{n=1}^9 \text{Normal}(x_n|0, \exp(y/2))$$

```
parameters {
  real y;
  vector[9] x;
}
model {
  y ~ normal(0, 3);
  x ~ normal(0, exp(y/2));
}
```



```
parameters {
  real y_raw;
  vector[9] x_raw;
}
transformed parameters {
  real y;
  vector[9] x;

  y = 3.0 * y_raw;
  x = exp(y/2) * x_raw;
}
model {
  y_raw ~ normal(0, 1);
  x_raw ~ normal(0, 1);
}
```



Stan Sampling Parameters

cognitive model

statistics

computing

parameter	description	constraint	default
iterations	number of MCMC samples (per chain)	int, > 0	2000
delta: δ	target Metropolis acceptance rate	$\delta \in [0, 1]$	0.80
stepsize: ε	initial HMC step size	real, $\varepsilon > 0$	2.0
max_treedepth: L	maximum HMC steps per iteration	int, $L > 0$	10

Typical adjustments

- Increase iterations
- Increase delta
- Decrease stepsize
- Might have to increase max_treedepth

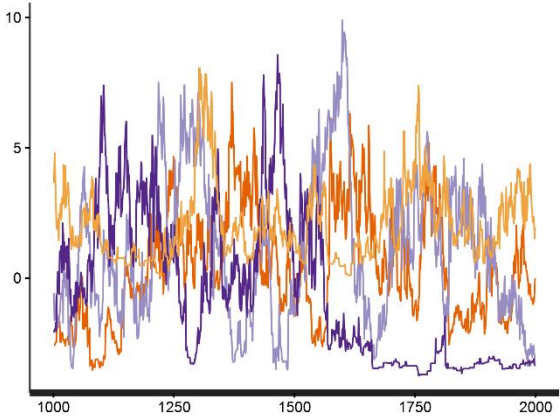
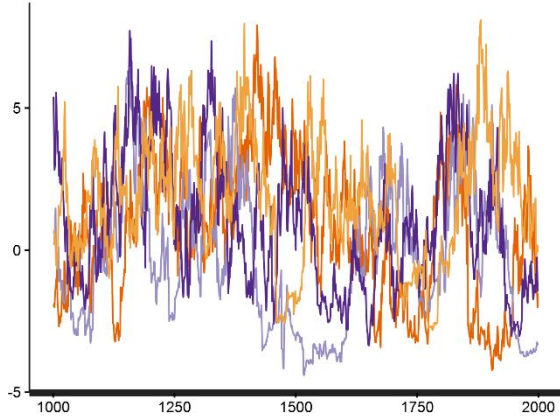
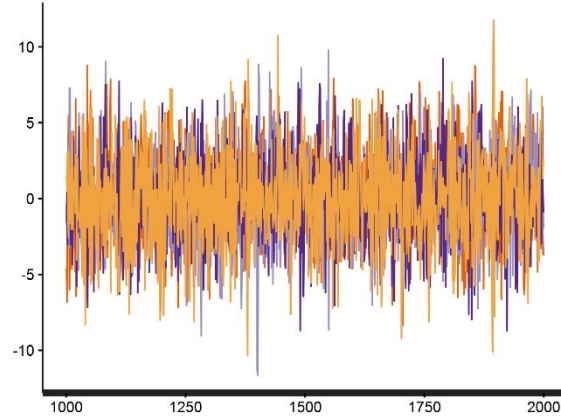
```
funnel_fit2 <- stan("_scripts/funnel.stan",  
  iter = 4000,  
  control = list(adapt_delta = 0.999,  
    stepsize = 1.0,  
    max_treedepth = 20))
```

Neal's Funnel: Comparing Performance

cognitive model

statistics

computing

	direct model	adjusted direct model	reparameterized model
Rhat (y)	1.22	1.1	1.0
n_eff (y)	18	42	3886
runtime*	48.50 sec	50.76 sec	50.12 sec
n_eff (y) / runtime	0.37 / sec	0.82 / sec	77.53 / sec
n_divergent	53	0	0
traceplot (y)			

*: 2 cores in parallel, including compiling time

How about **Bounded** Parameters?

$$\tilde{\theta} \sim \text{Normal}(0, 1)$$

$$\theta = \mu_{\theta} + \sigma_{\theta} \tilde{\theta}$$

$$\theta \in (-\infty, +\infty)$$



$$\tilde{\theta} \sim \text{Normal}(0, 1)$$

$$\theta = \text{Probit}^{-1}(\mu_{\theta} + \sigma_{\theta} \tilde{\theta})$$

$$\theta \in [0, 1]$$

constraint	reparameterization
$\theta \in (-\infty, +\infty)$	$\theta = \mu_{\theta} + \sigma_{\theta} \tilde{\theta}$
$\theta \in [0, N]$	$\theta = \text{Probit}^{-1}(\mu_{\theta} + \sigma_{\theta} \tilde{\theta}) \times N$
$\theta \in [M, N]$	$\theta = \text{Probit}^{-1}(\mu_{\theta} + \sigma_{\theta} \tilde{\theta}) \times (N-M) + M$
$\theta \in (0, +\infty)$	$\theta = \exp(\mu_{\theta} + \sigma_{\theta} \tilde{\theta})$

* Probit⁻¹: Normal cumulative distribution function (normcdf)

Apply to Our Hierarchical RL Model

cognitive model

statistics

computing

```
parameters {  
  real<lower=0,upper=1> lr_mu;  
  real<lower=0,upper=3> tau_mu;  
  
  real<lower=0> lr_sd;  
  real<lower=0> tau_sd;  
  
  real<lower=0,upper=1> lr[nSubjects];  
  real<lower=0,upper=3> tau[nSubjects];  
}
```



```
parameters {  
  # group-level parameters  
  real lr_mu_raw;  
  real tau_mu_raw;  
  real<lower=0> lr_sd_raw;  
  real<lower=0> tau_sd_raw;  
  
  # subject-level raw parameters  
  vector[nSubjects] lr_raw;  
  vector[nSubjects] tau_raw;  
}  
  
transformed parameters {  
  vector<lower=0,upper=1>[nSubjects] lr;  
  vector<lower=0,upper=3>[nSubjects] tau;  
  
  for (s in 1:nSubjects) {  
    lr[s] = Phi_approx( lr_mu_raw + lr_sd_raw * lr_raw[s] );  
    tau[s] = Phi_approx( tau_mu_raw + tau_sd_raw * tau_raw[s] ) * 3;  
  }  
}
```

Apply to Our Hierarchical RL Model

cognitive model

statistics

computing

```
model {  
  lr_sd ~ cauchy(0,1);  
  tau_sd ~ cauchy(0,3);  
  lr ~ normal(lr_mu, lr_sd) ;  
  tau ~ normal(tau_mu, tau_sd) ;  
  
  for (s in 1:nSubjects) {  
    vector[2] v;  
    real pe;  
    v = initV;  
  
    for (t in 1:nTrials) {  
      choice[s,t] ~ categorical_logit( tau[s] * v );  
      pe = reward[s,t] - v[choice[s,t]];  
      v[choice[s,t]] = v[choice[s,t]] + lr[s] * pe;  
    }  
  }  
}
```



```
model {  
  lr_mu_raw ~ normal(0,1);  
  tau_mu_raw ~ normal(0,1);  
  lr_sd_raw ~ cauchy(0,3);  
  tau_sd_raw ~ cauchy(0,3);  
  
  lr_raw ~ normal(0,1);  
  tau_raw ~ normal(0,1);  
  
  for (s in 1:nSubjects) {  
    ...  
  }  
  
  generated quantities {  
    real<lower=0,upper=1> lr_mu;  
    real<lower=0,upper=3> tau_mu;  
  
    lr_mu = Phi_approx(lr_mu_raw);  
    tau_mu = Phi_approx(tau_mu_raw) * 3;  
  }  
}
```

Exercise XII

cognitive model

statistics

computing

```
.../07.optm_rl/_scripts/reinforcement_learning_hrch_main.R
```

TASK: (1) Complete the Matt Trick
(2) fit the optimized hierarchical RL model

```
> source('_scripts/reinforcement_learning_hrch_main.R')  
> fit_rl4 <- run_rl_mp2(optimized = TRUE)
```

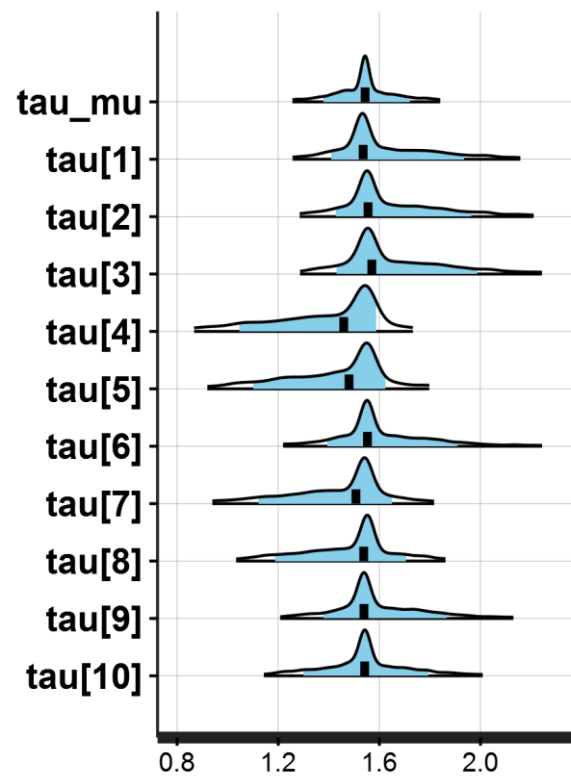
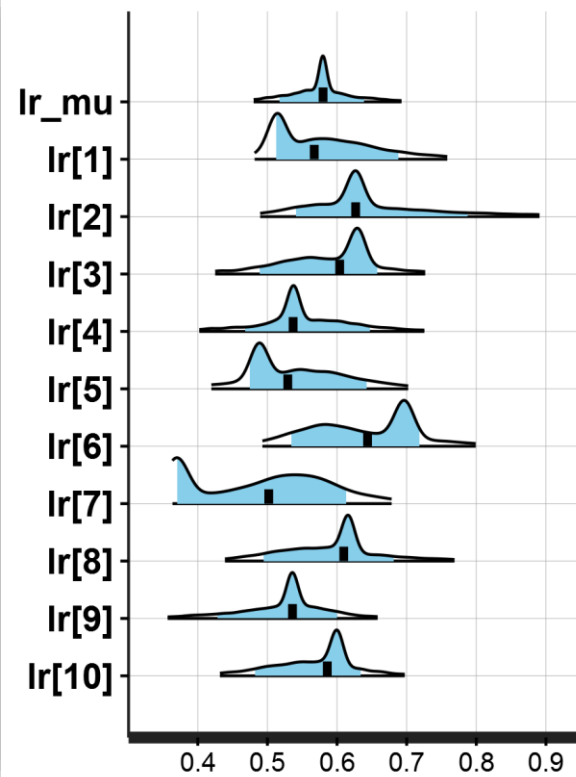
Hierarchical Fitting – Optimized

cognitive model

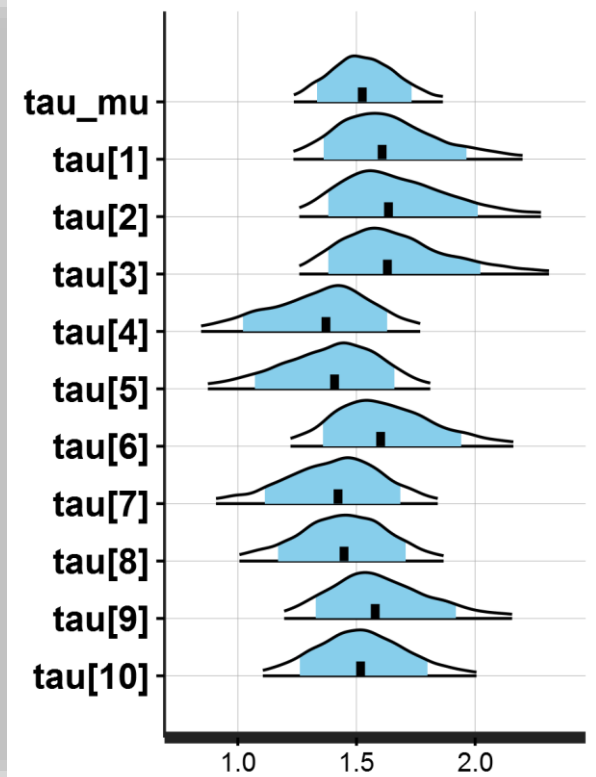
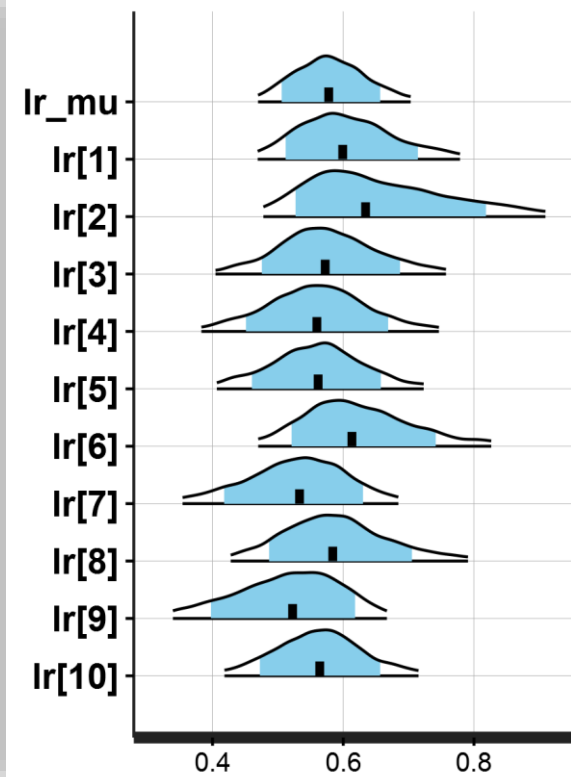
statistics

computing

Posterior Means (hrch)



Posterior Means (hrch + optm)



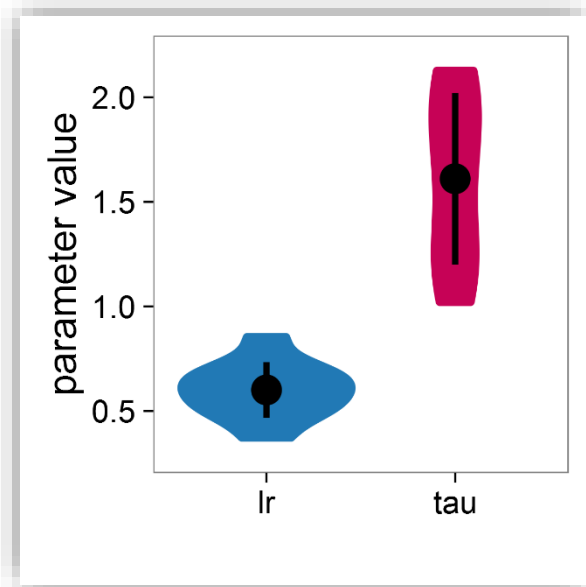
Comparing with True Parameters

cognitive model

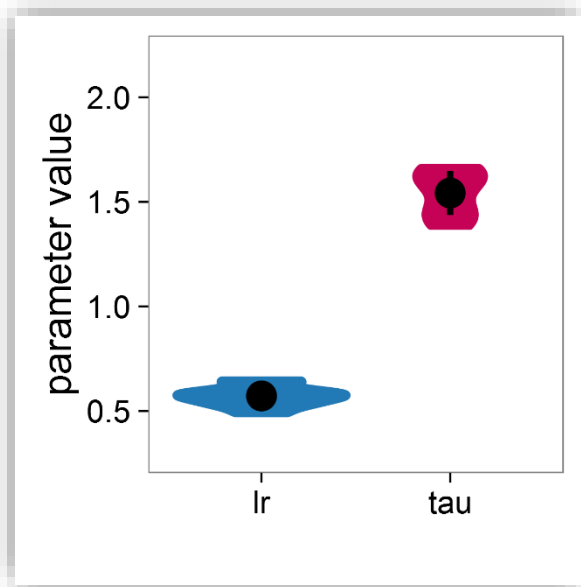
statistics

computing

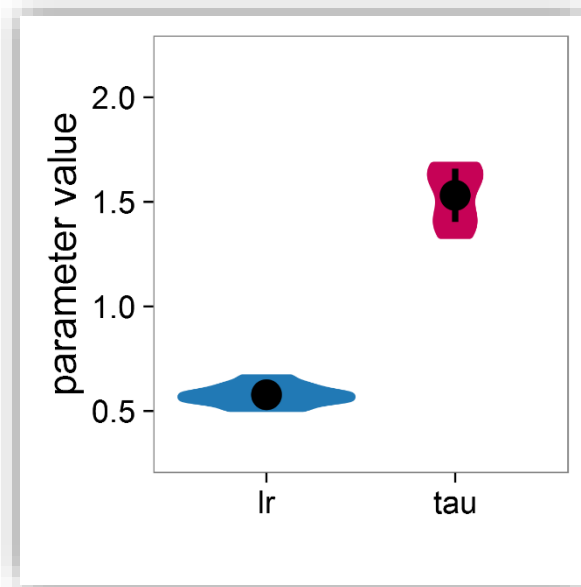
Posterior Means (indy)



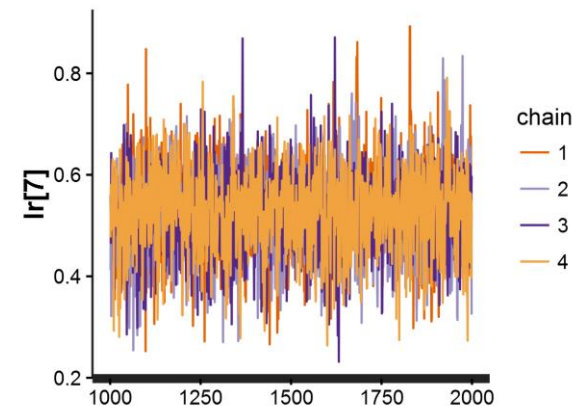
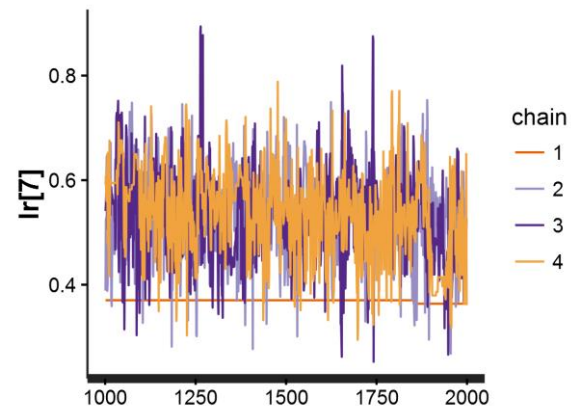
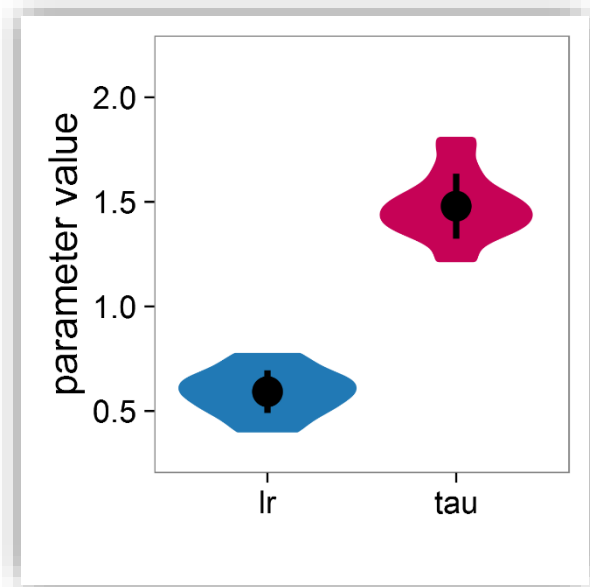
Posterior Means (hrch)



Posterior Means (hrch+optm)



True Parameters



Posterior Predictive Check

cognitive model

statistics

computing

